

MAY 15, 2025

Kowalski – Your Helpful AI Assistant

Christian Goll <cgoll@suse.com>



Outline

Introduction

AI Tools

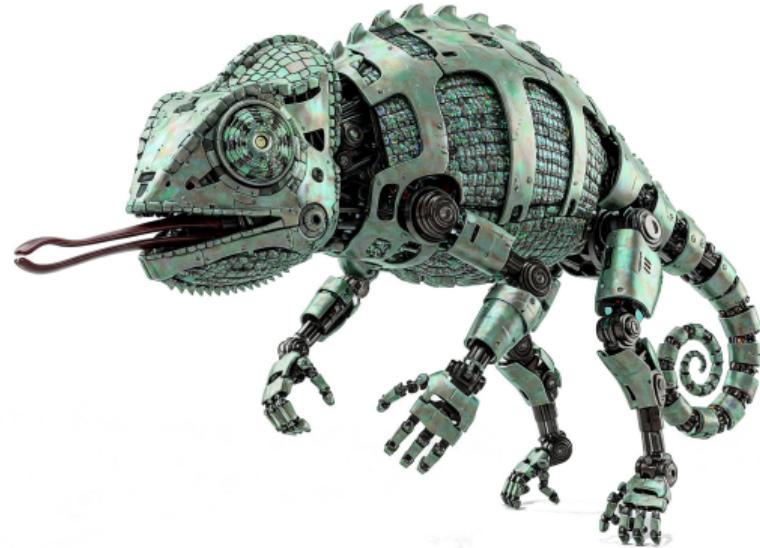
Kowalski Implementation



The Challenge: LLMs and System Configuration

Current state

- Large Language Models (LLMs) are powerful for general tasks.
- LLMs are already used for system configuration tasks
- Hallucinations and wrong tools likely to be presented



Example

- Answers with standard prompt are extensive
- User has to select the right one
- Sensitive information may be shared!



Great! Installing Vim on Linux is simple, and the method depends on your distribution. Here are the commands for the most common Linux package managers:

1. Debian/Ubuntu (and derivatives like Linux Mint, Pop!_OS)

```
bash Copy Download  
sudo apt update && sudo apt install vim -y
```

- * `sudo apt update` updates your package list.
- * `sudo apt install vim -y` installs Vim automatically without confirmation.

2. Fedora/RHEL/CentOS/AlmaLinux/Rocky Linux

```
bash Copy Download  
sudo dnf install vim      # Fedora (modern)  
sudo yum install vim      # Older RHEL/CentOS
```

3. Arch Linux (and Manjaro)

```
bash Copy Download  
sudo pacman -S vim
```

4. OpenSUSE

```
bash Copy Download  
sudo zypper install vim
```



Problem

- LLMs are created by training text completion on scraped sources
- LLMs don't have any context



You know nothing, John Snow!



Solution

- Context is needed for accurate solutions
- Should be easy to get context of actual system.



I know things!



Introducing Kowalski: A Specialized AI Assistant

- Kowalski is designed to bridge the gap between LLMs and system configuration.
- It provides LLMs with highly relevant and context-aware information.
- Focus on extracting and utilizing knowledge from technical documentation, specifically SUSE/SLE docs.



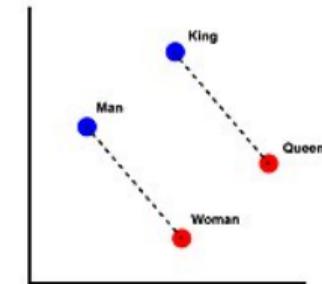
Vector Embeddings

Definition

Numerical representations of data (words, texts, etc.) in a vector space

Key Properties

- Capture semantic relationships
(e.g. "king" – "man" + "woman" \approx "queen").
- Fixed-length
- Similar items cluster together in the vector space.



Example

"cat" [0.2, -0.5, 0.7], "dog" [0.3, -0.4, 0.6]



Embedding details

Creation methods

GloVe

Uses global word co-occurrence statistics from a corpus.

Word2Vec

predicts surrounding words (or target word) using shallow neural networks.

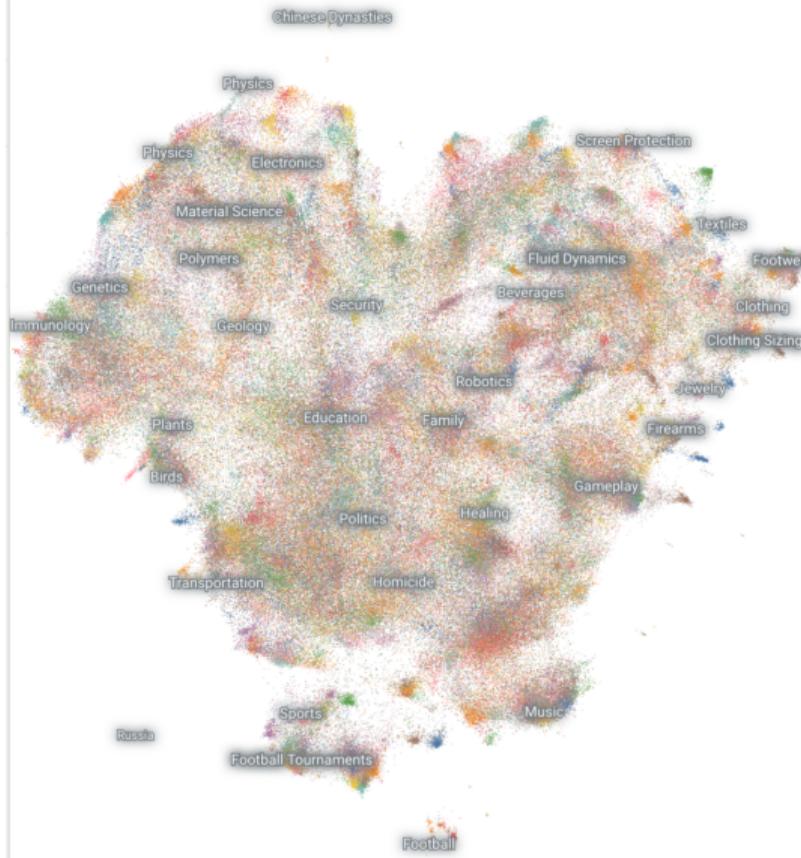
Vector space search

- Searching nearest neighbors in high dimensional vector space isn't trivial
- Facebook created `faiss` library which does this job



Embedding conclusions

- Use Word2Vec like embedding so that LLM and embedding run on Ollama
- Creation embedding for documentation is computational expensive
- Embeddings are available for with different OSS licenses, text lengths and embedding dimensions



Ollama: Local AI Powerhouse

- Ollams is Open-source framework for running LLMs locally
- Enables execution of models like LLaMA on consumer hardware

Efficiency

- Uses **quantization**—reduces model precision (e.g. 32-bit → 4-bit)
- Smaller model size + faster inference
- Minimal accuracy loss for significant performance gains

The LLaMA Foundation

- Meta's LLaMA was first major open weights model
- Proved smaller models could compete with proper training



Why Go for Kowalski?

Python's Challenges in Production AI

- **Dependency Hell**

- Python has a *long tail*—importing one package pulls in many others which have to be packaged
- Increases deployment complexity
- e.g. python-faiss is not available for the *oldest* python in Tumbleweed

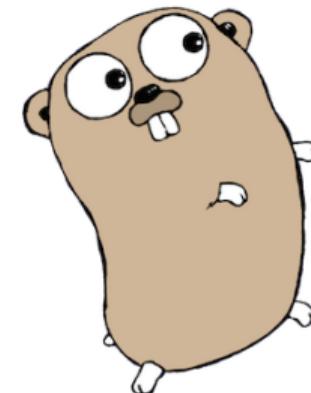
Go's Advantages

- **Single Binary Deployment**

- No dependency chains—statically linked executables

- **Compile Time Vendoring**

- Dependencies are imported at compile time and not run time



How does Kowalski work

